

```
In [7]: import numpy as np
import pandas as pd
train_df = pd.read_csv('train.csv', nrows=1000000)
train_df.shape
```

```
Out[7]: (1000000, 8)
```

```
In [8]: #remove unwanted data. After multiple iterations and runs, efficient way
to data clean

required_data = (train_df['fare_amount'].between(2.5, 200) & train_df['p
assenger_count'].between(0, 6) &
                train_df['pickup_longitude'].between(-74.5, -72.5) & tra
in_df['dropoff_longitude'].between(-74.5, -72.5) &
                train_df['pickup_latitude'].between(40, 42) & train_df[
'dropoff_latitude'].between(40, 42))
```

```
In [9]: train_df.shape
```

```
Out[9]: (1000000, 8)
```

```
In [10]: train_df = train_df[required_data]
```

```
In [11]: train_df.shape
```

```
Out[11]: (979060, 8)
```

```
In [16]: train_df.head()
```

```
Out[16]:
```

	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude
0	2009-06-15 17:26:21.0000001	4.5	2009-06-15 17:26:21 UTC	-73.844311	40.721319
1	2010-01-05 16:52:16.0000002	16.9	2010-01-05 16:52:16 UTC	-74.016048	40.711303
2	2011-08-18 00:35:00.00000049	5.7	2011-08-18 00:35:00 UTC	-73.982738	40.761270
3	2012-04-21 04:30:42.0000001	7.7	2012-04-21 04:30:42 UTC	-73.987130	40.733143
4	2010-03-09 07:51:00.000000135	5.3	2010-03-09 07:51:00 UTC	-73.968095	40.768008

```
train_df['pickup_datetime'] = pd.to_datetime(train_df['pickup_datetime'])
```

```
In [18]: train_df['pickup_datetime'] = pd.to_datetime(train_df['pickup_datetime']
])
```

```
In [19]: def extractdatetime(data):
    data['hour'] = data['pickup_datetime'].dt.hour
    data['year'] = data['pickup_datetime'].dt.year
    data['month'] = data['pickup_datetime'].dt.month
    data['date'] = data['pickup_datetime'].dt.day
    data['day'] = data['pickup_datetime'].dt.dayofweek
    return data

train_df=extractdatetime(train_df)
```

```
In [20]: train_df.head()
```

```
Out[20]:
```

	e_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
		2009-06-15 17:26:21	-73.844311	40.721319	-73.841610	40.71227
9		2010-01-05 16:52:16	-74.016048	40.711303	-73.979268	40.78200
		2011-08-18 00:35:00	-73.982738	40.761270	-73.991242	40.75056
		2012-04-21 04:30:42	-73.987130	40.733143	-73.991567	40.75809
		2010-03-09 07:51:00	-73.968095	40.768008	-73.956655	40.78376

```
In [145]: #Thanks to Stack_Overflow. Another efficient way of calculating distance
           when co-ordinates are given (Haversine Distance)
def distance(pickup_longitude, pickup_latitude, dropoff_longitude, dropoff_latitude):
    pickup_longitude, pickup_latitude, dropoff_longitude, dropoff_latitude = map(np.radians, [pickup_longitude, pickup_latitude, dropoff_longitude, dropoff_latitude])
    dlong = dropoff_longitude - pickup_longitude
    dlat = dropoff_latitude - pickup_latitude
    a = np.sin(dlat/2.0)**2 + np.cos(pickup_latitude) * np.cos(dropoff_latitude) * np.sin(dlong/2.0)**2
    c = 2 * np.arcsin(np.sqrt(a))
    distance = 6367 * c
    return distance

train_df['distance'] = distance(train_df['pickup_longitude'], train_df['pickup_latitude'],
                               train_df['dropoff_longitude'], train_df['dropoff_latitude'])
```

In [23]: `train_df.head(1)`

Out[23]:

	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dr
0	2009-06-15 17:26:21.0000001	4.5	2009-06-15 17:26:21	-73.844311	40.721319	-7

In [24]: `print(train_df.corrwith(train_df['fare_amount']))`

```

fare_amount          1.000000
pickup_longitude      0.380484
pickup_latitude     -0.189488
dropoff_longitude     0.290090
dropoff_latitude    -0.158196
passenger_count      0.014383
hour                 -0.019302
year                  0.117216
month                 0.025561
date                  0.001324
day                   0.002705
distance              0.821817
dtype: float64

```

```

In [25]: def absolute_coordinates(data):
          data['absolute_longitude'] = (data.dropoff_longitude - data.pickup_l
ongitude).abs()
          data['absolute_latitude'] = (data.dropoff_latitude - data.pickup_lat
itude).abs()
          return data

train_df=absolute_coordinates(train_df)

```

In [26]: `train_df.head(2)`

Out[26]:

	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dr
0	2009-06-15 17:26:21.0000001	4.5	2009-06-15 17:26:21	-73.844311	40.721319	-7
1	2010-01-05 16:52:16.0000002	16.9	2010-01-05 16:52:16	-74.016048	40.711303	-7

```

In [27]: def E_distance(lat1, long1, lat2, long2):
          sqlat=(lat1-lat2)**2
          sqlong=(long1-long2)**2
          e_distance = np.sqrt(sqlat+sqlong)
          return e_distance

```

```

In [28]: train_df['e_distance'] = E_distance(train_df.pickup_latitude, train_df.p
ickup_longitude,
                                             train_df.dropoff_latitude, train_df.dropo
ff_longitude)

```

```
In [29]: train_df.head(5)
```

```
Out[29]:
```

	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude
0	2009-06-15 17:26:21.0000001	4.5	2009-06-15 17:26:21	-73.844311	40.721319
1	2010-01-05 16:52:16.0000002	16.9	2010-01-05 16:52:16	-74.016048	40.711303
2	2011-08-18 00:35:00.00000049	5.7	2011-08-18 00:35:00	-73.982738	40.761270
3	2012-04-21 04:30:42.0000001	7.7	2012-04-21 04:30:42	-73.987130	40.733143
4	2010-03-09 07:51:00.000000135	5.3	2010-03-09 07:51:00	-73.968095	40.768008

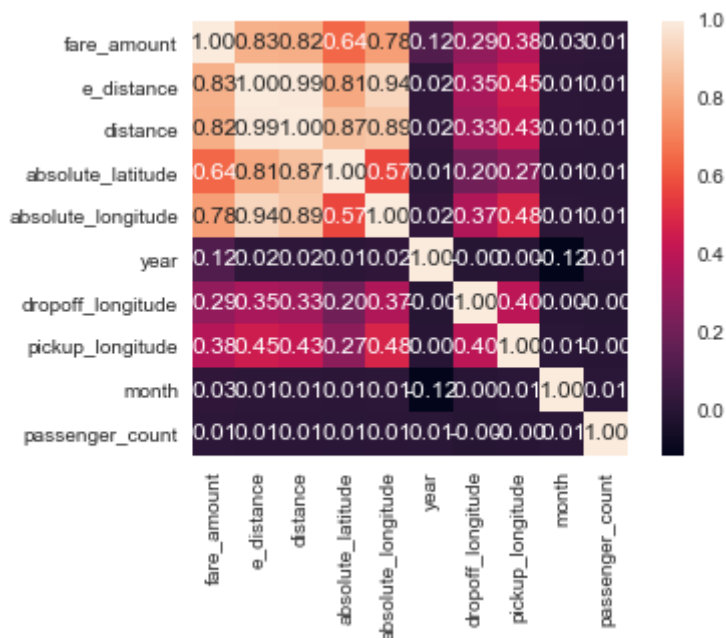
```
In [30]: print(train_df.corrwith(train_df['fare_amount']))
```

```
fare_amount          1.000000
pickup_longitude      0.380484
pickup_latitude     -0.189488
dropoff_longitude     0.290090
dropoff_latitude     -0.158196
passenger_count       0.014383
hour                 -0.019302
year                  0.117216
month                 0.025561
date                  0.001324
day                   0.002705
distance              0.821817
absolute_longitude    0.778014
absolute_latitude     0.643646
e_distance            0.830235
dtype: float64
```

```
In [146]: import seaborn as sns
import matplotlib.pyplot as plt

spearman_correlation = train_df.corr(method='spearman')
pick_columns=spearman_correlation.nlargest(10, 'fare_amount').index
correlationmap = np.corrcoef(train_df[pick_columns].values.T)
sns.set(font_scale=1.0)
heatmap = sns.heatmap(correlationmap, cbar=True, annot=True, square=True,
, fmt='.2f', yticklabels=columns_new.values, xticklabels=columns_new.val
ues)

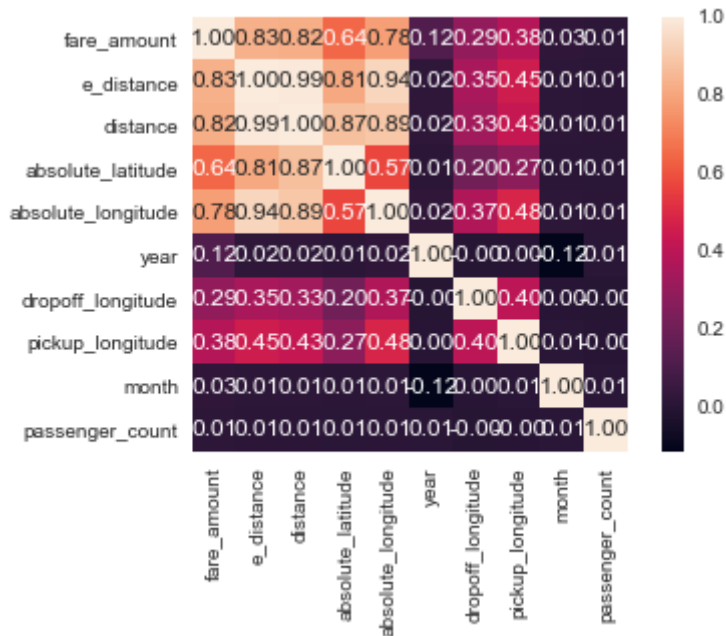
plt.show()
```



```
In [147]: import seaborn as sns
import matplotlib.pyplot as plt

spearman_correlation = train_df.corr(method='spearman')
pick_columns=spearman_correlation.nlargest(10, 'fare_amount').index
correlationmap = np.corrcoef(train_df[pick_columns].values.T)
sns.set(font_scale=1.0)
heatmap = sns.heatmap(correlationmap, cbar=True, annot=True, square=True,
, fmt='.2f', yticklabels=columns_new.values, xticklabels=columns_new.val
ues)

plt.show()
```



```
In [33]: print(train_df.corrwith(train_df['fare_amount']))
```

```
fare_amount      1.000000
pickup_longitude  0.380484
pickup_latitude -0.189488
dropoff_longitude 0.290090
dropoff_latitude -0.158196
passenger_count  0.014383
hour             -0.019302
year              0.117216
month             0.025561
date              0.001324
day               0.002705
distance          0.821817
absolute_longitude 0.778014
absolute_latitude  0.643646
e_distance        0.830235
dtype: float64
```

```
In [35]: train_df.corr(method='pearson')
```


```
Out[35]:
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	d
fare_amount	1.000000	0.380484	-0.189488	0.290090	-C
pickup_longitude	0.380484	1.000000	0.127059	0.400761	0.
pickup_latitude	-0.189488	0.127059	1.000000	0.152140	0.
dropoff_longitude	0.290090	0.400761	0.152140	1.000000	0.
dropoff_latitude	-0.158196	0.134755	0.477618	0.224905	1.
passenger_count	0.014383	-0.000187	-0.008182	-0.001703	-C
hour	-0.019302	0.017943	0.027865	-0.041459	0.
year	0.117216	0.002308	-0.019389	-0.000607	-C
month	0.025561	0.006070	-0.003308	0.003969	-C
date	0.001324	-0.000718	-0.001333	0.001591	-C
day	0.002705	-0.023948	-0.037401	-0.001651	-C
distance	0.821817	0.427645	-0.150194	0.331222	-C
absolute_longitude	0.778014	0.480205	-0.142469	0.370528	-C
absolute_latitude	0.643646	0.274334	-0.121836	0.204497	-C
e_distance	0.830235	0.449165	-0.151735	0.347684	-C

```
In [38]: import matplotlib.pyplot as plt

plt.scatter(train_df.distance, train_df.fare_amount)
plt.xlabel('distance mile')
plt.ylabel('fare price')


# theta here is estimated by hand
theta = (16, 4.0)
x = np.linspace(0.1, 3, 50)
plt.plot(x, theta[0]/x + theta[1], '--', c='r', lw=2);
```



```
In [39]: import matplotlib.pyplot as plt


plt.scatter(train_df.e_distance, train_df.fare_amount)
plt.xlabel('Euclidian distance mile')
plt.ylabel('fare price')

# theta here is estimated by hand
theta = (16, 4.0)
x = np.linspace(0.1, 3, 50)
plt.plot(x, theta[0]/x + theta[1], '--', c='r', lw=2);
```





```
In [41]: plt.figure(figsize=(15,7))  
plt.scatter(x=train_df['e_distance'], y=train_df['fare_amount'], s=1.5)  
plt.xlabel('euclidian distance')  
plt.ylabel('Fare')
```

Out[41]: Text(0,0.5,'Fare')




```
In [42]: plt.figure(figsize=(15,7))  
plt.scatter(x=train_df['distance'], y=train_df['fare_amount'], s=1.5)  
plt.xlabel('distance')  
plt.ylabel('Fare')
```

Out[42]: Text(0,0.5,'Fare')



```
In [43]: plt.figure(figsize=(15,7))
plt.scatter(x=train_df['hour'], y=train_df['fare_amount'], s=1.5)
plt.xlabel('time of a day')
plt.ylabel('Fare')
```


Out[43]: Text(0,0.5,'Fare')



```
In [44]: import matplotlib.pyplot as plt


plt.scatter(train_df.hour, train_df.fare_amount)
plt.xlabel('Time of a day')
plt.ylabel('fare price')

# theta here is estimated by hand
theta = (16, 4.0)
x = np.linspace(0.1, 3, 50)
plt.plot(x, theta[0]/x + theta[1], '--', c='r', lw=2);
```



```
In [45]: plt.figure(figsize=(15,7))
plt.scatter(x=train_df['hour'], y=train_df['e_distance'], s=2)
plt.xlabel('time of a day')
plt.ylabel('Eucledian distance')
```


Out[45]: Text(0,0.5,'Eucledian distance')



```
In [46]: import matplotlib.pyplot as plt

plt.scatter(train_df.hour, train_df.e_distance)
plt.xlabel('Time of a day')
plt.ylabel('eucledian distance')

# theta here is estimated by hand
theta = (16, 4.0)
x = np.linspace(0.1, 3, 50)
plt.plot(x, theta[0]/x + theta[1], '--', c='r', lw=2);
```




```
In [47]: train_df.head(2)
```

```
Out[47]:
```

	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dr
0	2009-06-15 17:26:21.0000001	4.5	2009-06-15 17:26:21	-73.844311	40.721319	-7
1	2010-01-05 16:52:16.0000002	16.9	2010-01-05 16:52:16	-74.016048	40.711303	-7


```
In [48]: plt.figure(figsize=(15,7))
plt.scatter(x=train_df['fare_amount'], y=train_df['absolute_longitude'],
            s=2)
plt.xlabel('Fare amount')
plt.ylabel('Absolute Longitude')
```

```
Out[48]: Text(0,0.5,'Absolute Longitude')
```




```
In [49]: plt.figure(figsize=(15,7))  
plt.scatter(x=train_df['fare_amount'], y=train_df['absolute_latitude'],  
s=2)  
plt.xlabel('Fare amount')  
plt.ylabel('Absolute Latitude')
```

Out[49]: Text(0,0.5,'Absolute Latitude')




```
In [51]: plt.figure(figsize=(15,7))  
plt.hist(train_df['hour'], bins=100)  
plt.xlabel('Hour')  
plt.ylabel('Frequency')
```

Out[51]: Text(0,0.5,'Frequency')




```
In [52]: plt.figure(figsize=(15,7))
plt.scatter(x=train_df['absolute_latitude'], y=train_df['fare_amount'],
s=2)
plt.xlabel('absolute latitude')
plt.ylabel('fare amount')
```

Out[52]: Text(0,0.5,'fare amount')




```
In [53]: plt.figure(figsize=(15,7))
plt.scatter(x=train_df['absolute_latitude'], y=train_df['e_distance'], s
=2)
plt.xlabel('absolute latitude')
plt.ylabel('e_distance')
```

Out[53]: Text(0,0.5,'e_distance')




```
In [54]: plt.figure(figsize=(15,7))
plt.scatter(x=train_df['absolute_longitude'], y=train_df['e_distance'],
s=2)
plt.xlabel('absolute longitude')
plt.ylabel('e_distance')
```

Out[54]: Text(0,0.5,'e_distance')




```
In [55]: plt.figure(figsize=(15,7))
plt.scatter(x=train_df['day'], y=train_df['e_distance'], s=2)
plt.xlabel('Week Day')
plt.ylabel('e_distance')
```

Out[55]: Text(0,0.5,'e_distance')




```
In [57]: plt.figure(figsize=(15,7))  
plt.hist(train_df['hour'], bins=100)  
plt.xlabel('Hour')  
plt.ylabel('Frequency')
```

Out[57]: Text(0,0.5,'Frequency')




```
In [59]: plt.figure(figsize=(15,7))  
plt.scatter(x=train_df['day'], y=train_df['fare_amount'], s=1.5)  
plt.xlabel('Day of Week')  
plt.ylabel('Fare')
```

Out[59]: Text(0,0.5,'Fare')





```
In [61]: plt.figure(figsize=(15,7))  
plt.scatter(x=train_df['date'], y=train_df['fare_amount'], s=1.5)  
plt.xlabel('Date')  
plt.ylabel('Fare')
```

Out[61]: Text(0,0.5,'Fare')




```
In [65]: plt.figure(figsize=(15,7))  
plt.hist(train_df['fare_amount'], bins=100)  
plt.xlabel('Fare')  
plt.ylabel('Frequency')
```

Out[65]: Text(0,0.5,'Frequency')



```
In [66]: plt.figure(figsize=(15,7))
plt.hist(train_df['e_distance'], bins=100)
plt.xlabel('e_distance')
plt.ylabel('Frequency')
```

```
Out[66]: Text(0,0.5,'Frequency')
```



```
In [67]: train_df.corr(method='pearson')
```

```
Out[67]:
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	d
fare_amount	1.000000	0.380484	-0.189488	0.290090	-0.158196
pickup_longitude	0.380484	1.000000	0.127059	0.400761	0.134755
pickup_latitude	-0.189488	0.127059	1.000000	0.152140	0.477618
dropoff_longitude	0.290090	0.400761	0.152140	1.000000	0.224905
dropoff_latitude	-0.158196	0.134755	0.477618	0.224905	1.000000
passenger_count	0.014383	-0.000187	-0.008182	-0.001703	-0.001703
hour	-0.019302	0.017943	0.027865	-0.041459	0.027865
year	0.117216	0.002308	-0.019389	-0.000607	-0.019389
month	0.025561	0.006070	-0.003308	0.003969	-0.003308
date	0.001324	-0.000718	-0.001333	0.001591	-0.001333
day	0.002705	-0.023948	-0.037401	-0.001651	-0.037401
distance	0.821817	0.427645	-0.150194	0.331222	-0.150194
absolute_longitude	0.778014	0.480205	-0.142469	0.370528	-0.142469
absolute_latitude	0.643646	0.274334	-0.121836	0.204497	-0.121836
e_distance	0.830235	0.449165	-0.151735	0.347684	-0.151735

```
In [68]: train_df.head(2)
```

```
Out[68]:
```

	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dr
0	2009-06-15 17:26:21.0000001	4.5	2009-06-15 17:26:21	-73.844311	40.721319	-7
1	2010-01-05 16:52:16.0000002	16.9	2010-01-05 16:52:16	-74.016048	40.711303	-7

```
In [69]: train_df.shape
```

```
Out[69]: (979060, 17)
```

```
In [70]: train_df_train = train_df[:600000]
train_df_test = train_df[600001:]
```

```
In [72]: from sklearn.pipeline import Pipeline
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error

train_df_train=train_df_train.drop(['key','pickup_datetime'], axis = 1)
train_df_test=train_df_test.drop(['key','pickup_datetime'], axis = 1)
```

```
In [73]: train_df_train.shape
```

```
Out[73]: (600000, 17)
```

```
In [74]: train_df_test.shape
```

```
Out[74]: (379059, 15)
```

```
In [75]: train_df_train=train_df_train.drop(['key','pickup_datetime'], axis = 1)
```

```
In [76]: train_df_train.shape
```

```
Out[76]: (600000, 15)
```

```
In [78]: train_df_test.head(1)
```

```
Out[78]:
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
612897	10.5	-74.005414	40.736925	-73.973531	40.754818

```
In [79]: train_df_train = train_df[:600000]
train_df_train=train_df_train.drop(['key','pickup_datetime'], axis = 1)
train_df_train.shape
```

```
Out[79]: (600000, 15)
```

```
In [80]: train_df_test.shape
```

```
Out[80]: (379059, 15)
```

```
In [81]: train_df_test.head(1)
```

```
Out[81]:
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
612897	10.5	-74.005414	40.736925	-73.973531	40.754818

```
In [82]: train_df_train.head(1)
```

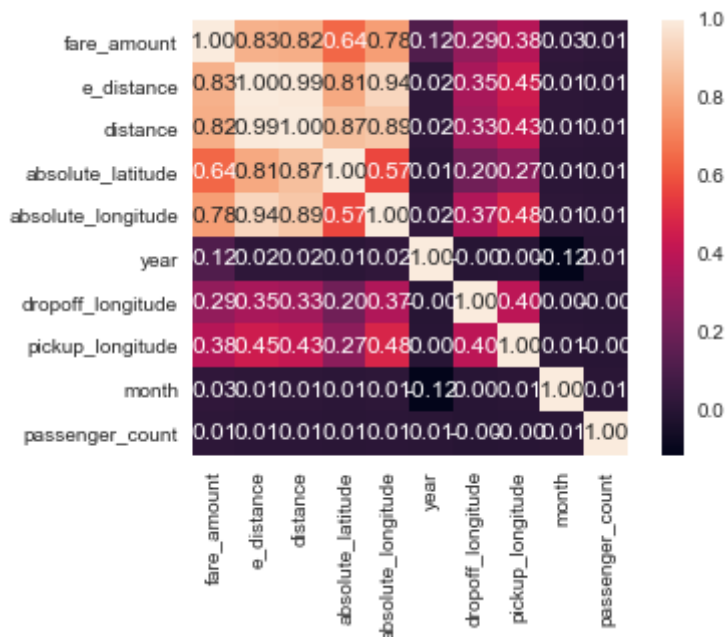
```
Out[82]:
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
0	4.5	-73.844311	40.721319	-73.84161	40.712278	1

```
In [148]: import seaborn as sns
import matplotlib.pyplot as plt

spearman_correlation = train_df.corr(method='spearman')
pick_columns=spearman_correlation.nlargest(10, 'fare_amount').index
correlationmap = np.corrcoef(train_df[pick_columns].values.T)
sns.set(font_scale=1.0)
heatmap = sns.heatmap(correlationmap, cbar=True, annot=True, square=True,
    , fmt='.2f', yticklabels=pick_columns_new.values, xticklabels=pick_columns_new.values)

plt.show()
```



```
In [84]: train_df_train.corr(method='pearson')
```

```
Out[84]:
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
fare_amount	1.000000	0.378509	-0.190201	0.293431	-0.160454
pickup_longitude	0.378509	1.000000	0.121854	0.399245	0.129408
pickup_latitude	-0.190201	0.121854	1.000000	0.146242	0.475894
dropoff_longitude	0.293431	0.399245	0.146242	1.000000	0.221165
dropoff_latitude	-0.160454	0.129408	0.475894	0.221165	1.000000
passenger_count	0.015512	0.000417	-0.007644	-0.002148	-0.007644
hour	-0.019601	0.016823	0.028402	-0.041793	0.028402
year	0.118150	0.002458	-0.019085	-0.000658	-0.019085
month	0.025915	0.005252	-0.004989	0.003540	-0.004989
date	0.001582	-0.000147	-0.001851	0.002497	-0.001851
day	0.003794	-0.023110	-0.037026	-0.000431	-0.037026
distance	0.820697	0.426361	-0.148876	0.336094	-0.148876
absolute_longitude	0.776904	0.477614	-0.141830	0.375410	-0.141830
absolute_latitude	0.642137	0.274753	-0.120143	0.207327	-0.120143
e_distance	0.829034	0.447487	-0.150511	0.352687	-0.150511

```
In [85]: train_df_train=train_df_train.drop(['hour','passenger_count','month',
        'date','day'], axis = 1)
        train_df_test=train_df_test.drop(['hour','passenger_count','month','date',
        'day'], axis = 1)
```

```
In [86]: train_df_train.head(1)
```

```
Out[86]:
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	year
0	4.5	-73.844311	40.721319	-73.84161	40.712278	2015

```
In [87]: train_df_test.head(1)
```

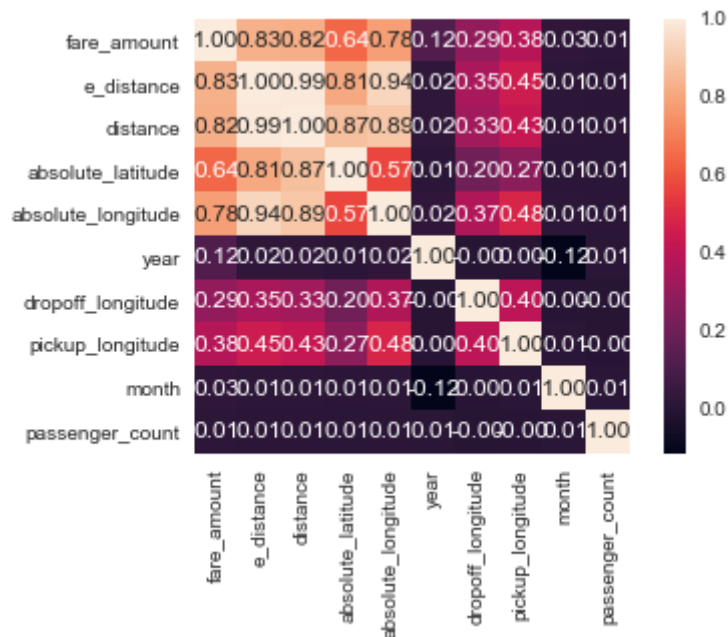
```
Out[87]:
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
612897	10.5	-74.005414	40.736925	-73.973531	40.754818

```
In [149]: import seaborn as sns
import matplotlib.pyplot as plt

spearman_correlation = train_df.corr(method='spearman')
pick_columns=spearman_correlation.nlargest(10, 'fare_amount').index
correlationmap = np.corrcoef(train_df[pick_columns].values.T)
sns.set(font_scale=1.0)
heatmap = sns.heatmap(correlationmap, cbar=True, annot=True, square=True,
fmt='.2f', yticklabels=columns_new.values, xticklabels=columns_new.values)

plt.show()
```



```
In [89]: X = train_df_train.drop('fare_amount',axis=1)
y = y = train_df_train[['fare_amount']]
```

```
In [91]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
random_state=42)
```

```
In [92]: from sklearn.metrics import mean_squared_error, explained_variance_score
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
```

```
In [150]: #Thanks to Abhishek Reddy Y N for helping me with this API
from sklearn.metrics import mean_squared_error

standard_scaler = StandardScaler().fit(X_train)
rescaled_X_train = standard_scaler.transform(X_train)
lin_model = LinearRegression()
lin_model.fit(rescaled_X_train, y_train)
pred = lin_model.predict(X_test)
error = np.sqrt(mean_squared_error(y_test, pred))
print(error)

1835.9237230723259
```

```
In [151]: print(pred)

[[1849.90727538]
 [1842.94903849]
 [1848.60092089]
 ...
 [1846.83974631]
 [1844.98424555]
 [1846.42364865]]
```

```
In [95]: train_df_train = train_df
```

```
In [96]: from sklearn.pipeline import Pipeline
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error

train_df_train=train_df_train.drop(['key', 'pickup_datetime'], axis = 1)
```

```
In [97]: train_df_train.head()
```


```
Out[97]:
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
0	4.5	-73.844311	40.721319	-73.841610	40.712278	1
1	16.9	-74.016048	40.711303	-73.979268	40.782004	1
2	5.7	-73.982738	40.761270	-73.991242	40.750562	2
3	7.7	-73.987130	40.733143	-73.991567	40.758092	1
4	5.3	-73.968095	40.768008	-73.956655	40.783762	1

```
In [152]: import seaborn as sns
import matplotlib.pyplot as plt

spearman_correlation = train_df.corr(method='spearman')
pick_columns=spearman_correlation.nlargest(10, 'fare_amount').index
correlationmap = np.corrcoef(train_df[pick_columns].values.T)
sns.set(font_scale=1.0)
heatmap = sns.heatmap(correlationmap, cbar=True, annot=True, square=True,
fmt='.2f', yticklabels=columns_new.values, xticklabels=columns_new.val
ues)

plt.show()
```



```
In [99]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
```

```
In [100]: from sklearn.metrics import mean_squared_error, explained_variance_score
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
```

```
In [153]: from sklearn.metrics import mean_squared_error

standard_scaler = StandardScaler().fit(X_train)
rescaled_X_train = standard_scaler.transform(X_train)
lin_model = LinearRegression()
lin_model.fit(rescaled_X_train, y_train)
pred = lin_model.predict(X_test)
error = np.sqrt(mean_squared_error(y_test, pred))
print(error)
```

1835.9237230723259

```
In [102]: train_df_train=train_df_train.drop(['passenger_count', 'hour', 'date', 'mon
th', 'day'], axis = 1)
```



```
In [103]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
```

```
In [104]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.5)
```

```
In [154]: from sklearn.metrics import mean_squared_error

standard_scaler = StandardScaler().fit(X_train)
rescaled_X_train = standard_scaler.transform(X_train)
lin_model = LinearRegression()
lin_model.fit(rescaled_X_train, y_train)
pred = lin_model.predict(X_test)
error = np.sqrt(mean_squared_error(y_test, pred))
print(error)
```

```
1835.9237230723259
```

```
In [141]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4)
```

```
In [155]: from sklearn.metrics import mean_squared_error

standard_scaler = StandardScaler().fit(X_train)
rescaled_X_train = standard_scaler.transform(X_train)
lin_model = LinearRegression()
lin_model.fit(rescaled_X_train, y_train)
pred = lin_model.predict(X_test)
error = np.sqrt(mean_squared_error(y_test, pred))
print(error)
```

```
1835.9237230723259
```

```
In [156]: print(pred)
```

```
[[1849.90727538]
 [1842.94903849]
 [1848.60092089]
 ...
 [1846.83974631]
 [1844.98424555]
 [1846.42364865]]
```

```
In [157]: train_df_train.corr(method='pearson')
```

```
Out[157]:
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
fare_amount	1.000000	0.380484	-0.189488	0.290090	-0.158196
pickup_longitude	0.380484	1.000000	0.127059	0.400761	0.134755
pickup_latitude	-0.189488	0.127059	1.000000	0.152140	0.477618
dropoff_longitude	0.290090	0.400761	0.152140	1.000000	0.224905
dropoff_latitude	-0.158196	0.134755	0.477618	0.224905	1.000000
year	0.117216	0.002308	-0.019389	-0.000607	-0.000607
distance	0.821817	0.427645	-0.150194	0.331222	-0.150194
absolute_longitude	0.778014	0.480205	-0.142469	0.370528	-0.142469
absolute_latitude	0.643646	0.274334	-0.121836	0.204497	-0.121836
e_distance	0.830235	0.449165	-0.151735	0.347684	-0.151735

```
In [158]: train_df_train.shape
```

```
Out[158]: (979060, 10)
```

```
In [159]: train_df_train.head(1)
```

```
Out[159]:
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	year
0	4.5	-73.844311	40.721319	-73.84161	40.712278	2010

```
In [160]: test_df = pd.read_csv('test.csv')
```

```
In [161]: test_df.shape
```

```
Out[161]: (9914, 7)
```

```
In [162]: test_df.isnull().sum().sort_values(ascending=False)
```

```
Out[162]: passenger_count    0
dropoff_latitude            0
dropoff_longitude           0
pickup_latitude             0
pickup_longitude            0
pickup_datetime             0
key                          0
dtype: int64
```

```
In [163]: test_df.head()
```

```
Out[163]:
```

	key	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude
0	2015-01-27 13:08:24.0000002	2015-01-27 13:08:24 UTC	-73.973320	40.763805	-73.981430
1	2015-01-27 13:08:24.0000003	2015-01-27 13:08:24 UTC	-73.986862	40.719383	-73.998886
2	2011-10-08 11:53:44.0000002	2011-10-08 11:53:44 UTC	-73.982524	40.751260	-73.979654
3	2012-12-01 21:12:12.0000002	2012-12-01 21:12:12 UTC	-73.981160	40.767807	-73.990448
4	2012-12-01 21:12:12.0000003	2012-12-01 21:12:12 UTC	-73.966046	40.789775	-73.988565

```
In [164]: test_df['pickup_datetime'] = pd.to_datetime(test_df['pickup_datetime'])
```

```
In [165]: test_df=extractdatetime(test_df)
test_df['distance'] = distance(test_df['pickup_longitude'], test_df['pickup_latitude'],
                               test_df['dropoff_longitude'], test_df['dropoff_latitude'])
test_df=absolute_coordinates(test_df)
test_df['e_distance'] = E_distance(test_df.pickup_latitude, test_df.pickup_longitude,
                                   test_df.dropoff_latitude, test_df.dropoff_longitude)
```

```
In [166]: test_df.head(2)
```

```
Out[166]:
```

	key	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude
0	2015-01-27 13:08:24.0000002	2015-01-27 13:08:24	-73.973320	40.763805	-73.981430
1	2015-01-27 13:08:24.0000003	2015-01-27 13:08:24	-73.986862	40.719383	-73.998886

```
In [167]: train_df_train.head(2)
```

```
Out[167]:
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	year
0	4.5	-73.844311	40.721319	-73.841610	40.712278	20
1	16.9	-74.016048	40.711303	-73.979268	40.782004	20

```
In [168]: test_df=test_df.drop(['key','pickup_datetime', 'hour', 'month', 'date',
'day'], axis = 1)
```

```
In [169]: train_df_train.head(2)
```

```
Out[169]:
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	year
0	4.5	-73.844311	40.721319	-73.841610	40.712278	2010
1	16.9	-74.016048	40.711303	-73.979268	40.782004	2010

```
In [170]: test_df.head(2)
```

```
Out[170]:
```

	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
0	-73.973320	40.763805	-73.981430	40.743835	1
1	-73.986862	40.719383	-73.998886	40.739201	1

```
In [173]: train_df_train.shape
```

```
Out[173]: (979060, 10)
```

```
In [172]: test_df.shape
```

```
Out[172]: (9914, 10)
```

```
In [174]: test_df=test_df.drop(['passenger_count'], axis = 1)
```

```
In [175]: test_df.shape
```

```
Out[175]: (9914, 9)
```

```
In [176]: train_df_train.shape
```

```
Out[176]: (979060, 10)
```

```
In [222]: #X = train_df_train.drop('fare_amount',axis=1)
#y = y = train_df_train[['fare_amount']]
X_test = test_df
print(len(X_test))
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)

standard_scaler = StandardScaler().fit(X_train)
rescaled_X_train = standard_scaler.transform(X_train)
lin_model = LinearRegression()
lin_model.fit(rescaled_X_train, y_train)
pred = lin_model.predict(X_test)
error = np.sqrt(mean_squared_error(y_test,pred))
print(error)
print (len(pred))

standard_scaler = StandardScaler().fit(y_train)
rescaled_y_train = standard_scaler.transform(y_train)

lin_model.fit(X, y)

9914
1844.7750999716266
180000
```

```
Out[222]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
```

```
In [226]: pred = lin_model.predict(test_df)
print (len(pred))
#error = np.sqrt(mean_squared_error(test_df,pred))

9914
```

```
In [227]: submission = pd.read_csv('sample_submission.csv')
submission['fare_amount'] = pred
submission.to_csv('submission_1.csv', index=False)
submission.head(20)
```

Out[227]:

	key	fare_amount
0	2015-01-27 13:08:24.0000002	10.405672
1	2015-01-27 13:08:24.0000003	11.167005
2	2011-10-08 11:53:44.0000002	5.430308
3	2012-12-01 21:12:12.0000002	7.862343
4	2012-12-01 21:12:12.0000003	13.630410
5	2012-12-01 21:12:12.0000005	10.292270
6	2011-10-06 12:10:20.0000001	6.284617
7	2011-10-06 12:10:20.0000003	50.174724
8	2011-10-06 12:10:20.0000002	11.589737
9	2014-02-18 15:22:20.0000002	7.770670
10	2014-02-18 15:22:20.0000003	10.199812
11	2014-02-18 15:22:20.0000001	15.172109
12	2010-03-29 20:20:32.0000002	5.345419
13	2010-03-29 20:20:32.0000001	7.997740
14	2011-10-06 03:59:12.0000002	9.142839
15	2011-10-06 03:59:12.0000001	12.939857
16	2012-07-15 16:45:04.0000006	5.849995
17	2012-07-15 16:45:04.0000002	9.079442
18	2012-07-15 16:45:04.0000003	6.163742
19	2012-07-15 16:45:04.0000004	5.938246

```
In [228]: submission.describe()
```

```
Out[228]:
```

	fare_amount
count	9914.000000
mean	11.556046
std	8.410039
min	-8.622536
25%	7.251999
50%	9.114863
75%	12.452505
max	180.442190

```
In [229]: import numpy as np
import pandas as pd
df = pd.read_csv('train.csv', nrows=5000000)
df.shape
```

```
Out[229]: (5000000, 8)
```

```
In [230]: required_data = (df['fare_amount'].between(2.5, 200) & df['passenger_cou
nt'].between(0, 6) &
df['pickup_longitude'].between(-74.5, -72.5) & df['dropo
ff_longitude'].between(-74.5, -72.5) &
df['pickup_latitude'].between(40, 42) & df['dropoff_lati
tude'].between(40, 42))
```

```
In [231]: df = df[required_data]
```

```
In [232]: df['pickup_datetime'] = pd.to_datetime(df['pickup_datetime'])
```

```
In [234]: def extractdatetime(data):
    data['hour'] = data['pickup_datetime'].dt.hour
    data['year'] = data['pickup_datetime'].dt.year
    data['month'] = data['pickup_datetime'].dt.month
    data['date'] = data['pickup_datetime'].dt.day
    data['day'] = data['pickup_datetime'].dt.dayofweek
    return data

df=extractdatetime(df)
```

```
In [235]: #Thanks to Stack_Overflow. Another efficient way of calculating distnce
when co-or are given (Haversine Distance)
def distance(pickup_longitude, pickup_latitude, dropoff_longitude, dropo
ff_latitude):
    pickup_longitude, pickup_latitude, dropoff_longitude, dropoff_latitu
de = map(np.radians, [pickup_longitude,pickup_latitude,dropoff_longitude
,dropoff_latitude])
    dlong = dropoff_longitude - pickup_longitude
    dlat = dropoff_latitude - pickup_latitude
    a = np.sin(dlat/2.0)**2 + np.cos(pickup_latitude) * np.cos(dropoff_l
atitude) * np.sin(dlong/2.0)**2
    c = 2 * np.arcsin(np.sqrt(a))
    distance = 6367 * c
    return distance

df['distance'] = distance(df['pickup_longitude'], df['pickup_latitude'],
                        df['dropoff_longitude'], df['dropoff_latitud
e'])
```

```
In [236]: def absolute_coordinates(data):
    data['absolute_longitude'] = (data.dropoff_longitude - data.pickup_l
ongitude).abs()
    data['absolute_latitude'] = (data.dropoff_latitude - data.pickup_lat
itude).abs()
    return data

df=absolute_coordinates(df)
```

```
In [237]: def E_distance(lat1, long1, lat2, long2):
    sqlat=(lat1-lat2)**2
    sqlong=(long1-long2)**2
    e_distance = np.sqrt(sqlat+sqlong)
    return e_distance

df['e_distance'] = E_distance(df.pickup_latitude, df.pickup_longitude,
                        df.dropoff_latitude, df.dropoff_longitude
)
```

```
In [238]: df.shape
```

```
Out[238]: (4893872, 17)
```

```
In [239]: df.head(2)
```

```
Out[239]:
```

	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dr
0	2009-06-15 17:26:21.0000001	4.5	2009-06-15 17:26:21	-73.844311	40.721319	-7
1	2010-01-05 16:52:16.0000002	16.9	2010-01-05 16:52:16	-74.016048	40.711303	-7

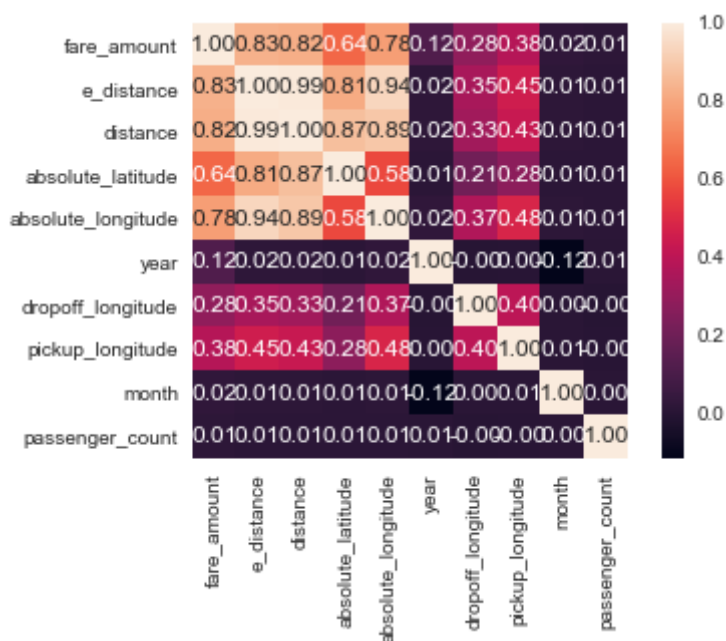

```
In [240]: print(df.corrwith(train_df['fare_amount']))
```

```
fare_amount      1.000000
pickup_longitude  0.380484
pickup_latitude  -0.189488
dropoff_longitude  0.290090
dropoff_latitude  -0.158196
passenger_count   0.014383
hour             -0.019302
year              0.117216
month             0.025561
date              0.001324
day               0.002705
distance          0.821817
absolute_longitude 0.778014
absolute_latitude  0.643646
e_distance        0.830235
dtype: float64
```

```
In [241]: import seaborn as sns
import matplotlib.pyplot as plt

spearman_correlation = df.corr(method='spearman')
pick_columns=spearman_correlation.nlargest(10, 'fare_amount').index
correlationmap = np.corrcoef(df[pick_columns].values.T)
sns.set(font_scale=1.0)
heatmap = sns.heatmap(correlationmap, cbar=True, annot=True, square=True,
    , fmt='.2f', yticklabels=pick_columns_new.values, xticklabels=pick_columns_new.val
ues)

plt.show()
```



```
In [243]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
```

```
In [252]: X = train_df_train.drop('fare_amount',axis=1)
          y = y = train_df_train[['fare_amount']]

          X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)

          standard_scaler = StandardScaler().fit(X_train)
          rescaled_X_train = standard_scaler.transform(X_train)

          from sklearn.ensemble import RandomForestRegressor
          rf = RandomForestRegressor()
          rf.fit(rescaled_X_train, y_train)
          rf_predict = rf.predict(X_test)
          #error = np.sqrt(mean_squared_error(y_test,pred))
          print(rf_predict)

/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:12: DataCo
nversionWarning: A column-vector y was passed when a 1d array was expec
ted. Please change the shape of y to (n_samples,), for example using ra
vel().
    if sys.path[0] == '':

[41.45 44.95 44.95 ... 44.95 44.95 41.45]
```

```
In [246]: train_df_train.shape
```

```
Out[246]: (979060, 10)
```

```
In [247]: len(y_test)
```

```
Out[247]: 293718
```

```
In [248]: len(X_test)
```

```
Out[248]: 293718
```

```
In [249]: len(rescaled_X_train)
```

```
Out[249]: 685342
```

```
In [250]: len(y_train)
```

```
Out[250]: 685342
```

```
In [253]: rf.fit(X, y)
```

```
/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:1: DataCon
versionWarning: A column-vector y was passed when a 1d array was expect
ed. Please change the shape of y to (n_samples,), for example using rav
el().
    """Entry point for launching an IPython kernel.
```

```
Out[253]: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,
                                max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=1,
                                oob_score=False, random_state=None, verbose=0, warm_start=False)
```

```
In [254]: rf_predict = rf.predict(X_test)
```

```
In [255]: print (rf_predict)
```

```
[20.2   3.73  5.35 ...  6.06  4.5   19.26]
```

```
In [257]: y_train.head(5)
```

```
Out[257]:
```

	fare_amount
661423	7.5
918011	16.9
960510	5.7
338825	7.0
991454	11.5

```
In [261]: for i in range(10):
            print (rf_predict[i])
```

```
20.2
3.7300000000000004
5.35
8.42
7.1400000000000001
10.74
10.27
21.883
4.5
9.6200000000000001
```

```
In [262]: y_train.head(10)
```

```
Out[262]:
```

	fare_amount
661423	7.5
918011	16.9
960510	5.7
338825	7.0
991454	11.5
954904	8.1
750031	19.3
475327	14.5
974243	7.7
1830	20.0

```
In [267]: rf_predict = rf.predict(X_test)
error = np.sqrt(mean_squared_error(y_test, rf_predict))
print(error)
```

```
1.6738955465124932
```

```
In [264]: print(rf_predict)
```

```
[ 8.3   10.2   5.06   ... 55.324 18.7    7.55 ]
```

```
In [268]: test_df.shape
```

```
Out[268]: (9914, 9)
```

```
In [269]: X.shape
```

```
Out[269]: (979060, 9)
```

```
In [270]: rf_predict = rf.predict(test_df)
print (rf_predict)
print(pred)
```

```
[ 8.3   10.2   5.06   ... 55.324 18.7    7.55 ]
[[10.40567211]
 [11.1670048 ]
 [ 5.43030812]
 ...
 [49.22548212]
 [20.60224933]
 [ 8.64386011]]
```

```
In [278]: from sklearn.linear_model import Ridge
ridge = Ridge(alpha=0.1)
ridge.fit(X_train, y_train)
ridge_predict_test = ridge.predict(X_test)
error = np.sqrt(mean_squared_error(y_test,ridge_predict_test))
print(error)
print(ridge_predict_test)

ridge.fit(X, y)
ridge_predict=ridge.predict(test_df)
print(ridge_predict)
```

```
5.125687978892424
[[17.61463391]
 [ 6.27831087]
 [ 6.51227483]
 ...
 [ 8.14029401]
 [ 5.602278   ]
 [17.59841975]]
[[10.41173641]
 [11.16332172]
 [ 5.42224178]
 ...
 [49.09186726]
 [20.72469083]
 [ 8.61227372]]
```

```
In [274]: submission['fare_amount'] = rf_predict  
submission.to_csv('submission_3.csv', index=False)  
submission.head(20)
```

Out[274]:

	key	fare_amount
0	2015-01-27 13:08:24.00000002	8.300
1	2015-01-27 13:08:24.00000003	10.200
2	2011-10-08 11:53:44.00000002	5.060
3	2012-12-01 21:12:12.00000002	7.650
4	2012-12-01 21:12:12.00000003	14.620
5	2012-12-01 21:12:12.00000005	9.800
6	2011-10-06 12:10:20.00000001	4.820
7	2011-10-06 12:10:20.00000003	48.360
8	2011-10-06 12:10:20.00000002	10.420
9	2014-02-18 15:22:20.00000002	6.150
10	2014-02-18 15:22:20.00000003	9.500
11	2014-02-18 15:22:20.00000001	15.400
12	2010-03-29 20:20:32.00000002	4.020
13	2010-03-29 20:20:32.00000001	7.940
14	2011-10-06 03:59:12.00000002	8.840
15	2011-10-06 03:59:12.00000001	19.454
16	2012-07-15 16:45:04.00000006	9.000
17	2012-07-15 16:45:04.00000002	8.640
18	2012-07-15 16:45:04.00000003	5.290
19	2012-07-15 16:45:04.00000004	4.650

```
In [279]: submission['fare_amount'] = ridge_predict
submission.to_csv('submission_4.csv', index=False)
submission.head(20)
```

Out[279]:

	key	fare_amount
0	2015-01-27 13:08:24.0000002	10.411736
1	2015-01-27 13:08:24.0000003	11.163322
2	2011-10-08 11:53:44.0000002	5.422242
3	2012-12-01 21:12:12.0000002	7.884502
4	2012-12-01 21:12:12.0000003	13.716213
5	2012-12-01 21:12:12.0000005	10.329417
6	2011-10-06 12:10:20.0000001	6.287167
7	2011-10-06 12:10:20.0000003	50.044504
8	2011-10-06 12:10:20.0000002	11.622250
9	2014-02-18 15:22:20.0000002	7.750008
10	2014-02-18 15:22:20.0000003	10.200876
11	2014-02-18 15:22:20.0000001	15.134557
12	2010-03-29 20:20:32.0000002	5.324779
13	2010-03-29 20:20:32.0000001	8.051190
14	2011-10-06 03:59:12.0000002	9.152499
15	2011-10-06 03:59:12.0000001	13.007124
16	2012-07-15 16:45:04.0000006	5.820099
17	2012-07-15 16:45:04.0000002	9.071017
18	2012-07-15 16:45:04.0000003	6.149045
19	2012-07-15 16:45:04.0000004	5.932387

```
In [281]: from sklearn.linear_model import Lasso
lassoreg = Lasso(alpha=0.1,normalize=True, max_iter=1e5)

lassoreg.fit(X_train, y_train)
lassoreg_predict_test = ridge.predict(X_test)
error_lasso = np.sqrt(mean_squared_error(y_test,lassoreg_predict_test))
print(error_lasso)
print(lassoreg_predict_test)

ridge.fit(X, y)
lassoreg_predict=ridge.predict(test_df)
print(lassoreg_predict)

5.125332979742834
[[17.59400102]
 [ 6.26921835]
 [ 6.51012099]
 ...
 [ 8.14375733]
 [ 5.59645646]
 [17.60968464]]
[[10.41173641]
 [11.16332172]
 [ 5.42224178]
 ...
 [49.09186726]
 [20.72469083]
 [ 8.61227372]]
```



```
In [282]: submission['fare_amount'] = lassoreg_predict
submission.to_csv('submission_5.csv', index=False)
submission.head(20)
```

Out[282]:

	key	fare_amount
0	2015-01-27 13:08:24.0000002	10.411736
1	2015-01-27 13:08:24.0000003	11.163322
2	2011-10-08 11:53:44.0000002	5.422242
3	2012-12-01 21:12:12.0000002	7.884502
4	2012-12-01 21:12:12.0000003	13.716213
5	2012-12-01 21:12:12.0000005	10.329417
6	2011-10-06 12:10:20.0000001	6.287167
7	2011-10-06 12:10:20.0000003	50.044504
8	2011-10-06 12:10:20.0000002	11.622250
9	2014-02-18 15:22:20.0000002	7.750008
10	2014-02-18 15:22:20.0000003	10.200876
11	2014-02-18 15:22:20.0000001	15.134557
12	2010-03-29 20:20:32.0000002	5.324779
13	2010-03-29 20:20:32.0000001	8.051190
14	2011-10-06 03:59:12.0000002	9.152499
15	2011-10-06 03:59:12.0000001	13.007124
16	2012-07-15 16:45:04.0000006	5.820099
17	2012-07-15 16:45:04.0000002	9.071017
18	2012-07-15 16:45:04.0000003	6.149045
19	2012-07-15 16:45:04.0000004	5.932387

```
In [283]: train_df.head(2)
```

Out[283]:

	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude
0	2009-06-15 17:26:21.0000001	4.5	2009-06-15 17:26:21	-73.844311	40.721319	-73.844311
1	2010-01-05 16:52:16.0000002	16.9	2010-01-05 16:52:16	-74.016048	40.711303	-74.016048

```
In [1]: from sklearn.neural_network import MLPClassifier
```

```
In [2]: mlp = MLPClassifier(hidden_layer_sizes=(13,13,13),max_iter=500)
```